

NLP Project

Full Circle

Vsevolod Dyomkin
prj-nlp-2, 2019-03-21



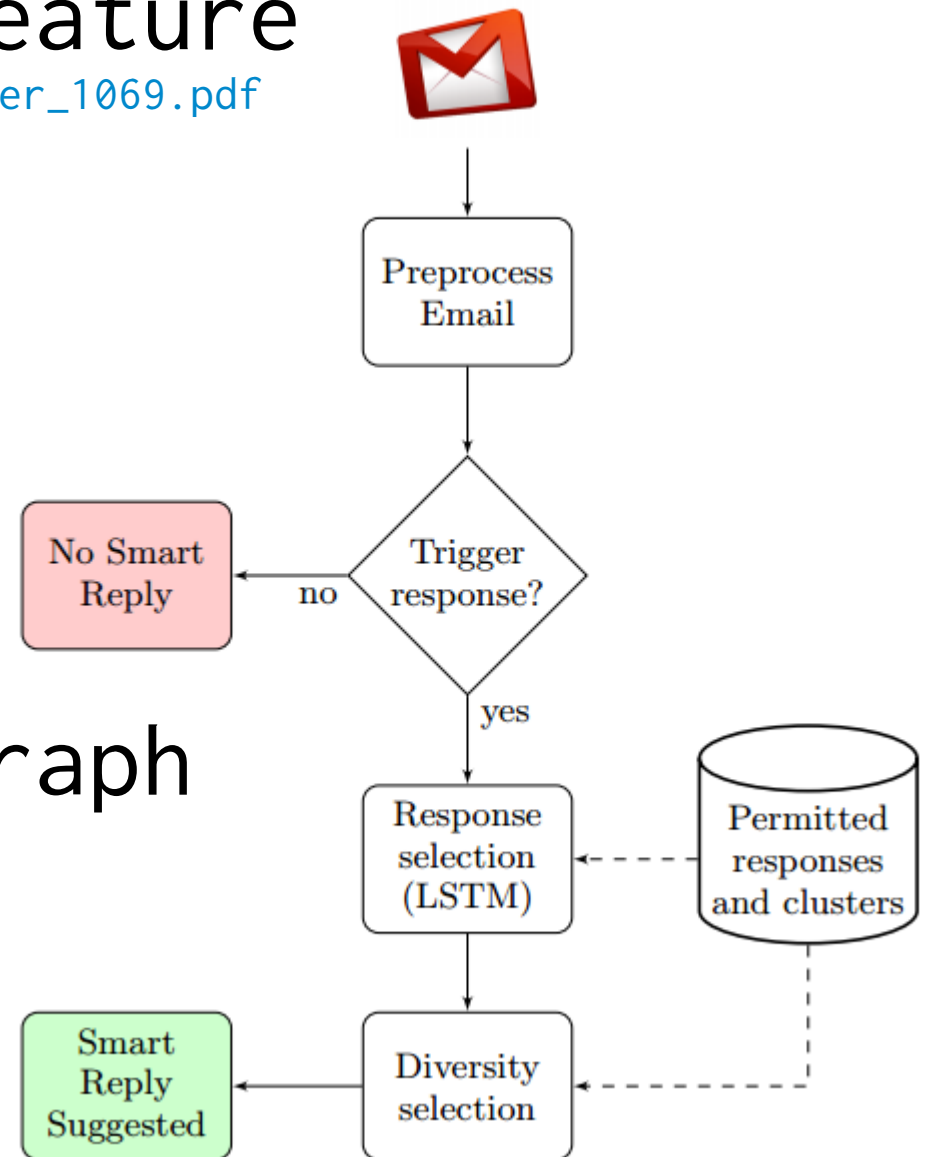
Глокая куздра штеко будланула бокра
И кудрячит бокрѣнка

Example

Gmail Smart Reply Feature

http://www.kdd.org/kdd2016/papers/files/Paper_1069.pdf


- * Data analysis
- * Traditional NLP processing
- * FNN
- * LSTM
- * Semi-supervised graph learning
- * Rule-based post-processing
- * Engineering



NLP Project Stages

- 1) Domain analysis
- 2) Data preparation
- 3) Iterating the solution
- 4) Productizing the result
- 5) Gathering feedback
and reiterating

Domain Analysis



Domain Analysis

- 1) Problem formulation
- 2) Existing solutions
- 3) Possible datasets
- 4) Identifying challenges
- 5) Metrics, baseline & SOTA
- 6) Selecting possible approaches and devising a plan of attack



Глокая куздра штеко будланула бокра
И кудрячит бокрѐнка

Problem Formulation

- * Not an unsolved problem
- * Short texts
- * Support “all” languages
- * The issue of UND/MULT texts
- * Dialects?
- * Encodings?

Interesting Examples

Text	Language	Explanation
Justin Bieber <3	und (Undefined)	NOT English; contains only a name.
Schalke XI v Chelsea: Fahrmann, Neustadter, Santana, Howedes, Uchida, Fuchs, Kirchhoff, Boateng, Hoger, Choupo-Moting, Huntelaar.	und (Undefined)	Contains only place/team/player names.
Ate spaghetti at La <u>tratoria napolitana</u>	en (English)	The name of the restaurant is in Italian, but the "main" language is English. An English-only speaker would understand this Tweet.
#NowListening Universo - Lodovica Comello @XYZ @XYZ	und (Undefined)	Italian song title and artist are just names. #NowListening is English but could be used by non-English speaker too.
#My #hot #naughty #neighbour #in #dallas: http://t.co/0dLJ 北京	en (English)	There is a Chinese word at the end, but the strongly prevailing language is English
Hahaha (•_•) (•_•)>~■-■ (~■_■) YEAHHH!	und (Undefined)	Emoticons and interjections only.
Que bonito!	und (Undefined)	Could be both Spanish and Portuguese
Pozor pozor	und (Undefined)	Could be Czech, Serbian, Croatian, Slovenian, ...
So warm in Berlin!	und (Undefined)	A valid sentence in both German and English
"Last Christmas" - <u>Der</u> Jose Carreras <u>unter den</u> <u>Weihnachtsliedern</u> .	de (German)	Contains an English song title and Spanish name, but is understandable to a German-only speaker.
Bécs <3	hu (Hungarian)	This is the Hungarian name for "Vienna", which is a proper name, but exists only in Hungarian
Estoy muy cansado voy a acostarme sooo tired <u>goi</u> n to <u>bedd</u>	und (Undefined)	Strong mixture of Spanish and English, no clear "main" language

<https://blog.twitter.com/2015/evaluating-language-identification-performance>

Existing Solutions


- * <https://github.com/shuyo/language-detection/> (Java)
- * <https://github.com/saffsd/langid.py> (Python)
- * <https://github.com/mzsanford/cld> (C++)

Existing Solutions

- * <https://github.com/shuyo/language-detection/> (Java)
- * <https://github.com/saffsd/langid.py> (Python)
- * <https://github.com/mzsanford/cld> (C++)
- * <https://github.com/CLD2Owners/cld2>
- * <https://github.com/google/cld3>

Existing Solutions

- * <https://github.com/shuyo/language-detection/> (Java)
- * <https://github.com/saffsd/langid.py> (Python)
- * <https://github.com/mzsanford/cld> (C++)
- * <https://github.com/CLD2Owners/cld2>
- * <https://github.com/google/cld3>



Possible Datasets

- * Debian i18n (~90 langs)

Possible Datasets

- * Debian i18n (~90 langs)
- * TED Multilingual (109 langs)

Possible Datasets

- * Debian i18n (~90 langs)
- * TED Multilingual (109 langs)
- * Wiktionary (~150 langs)

Possible Datasets

- * Debian i18n (~90 langs)
- * TED Multilingual (109 langs)
- * Wiktionary (~150 langs)
- * Wikipedia (175 langs)

Test Data

- * Twitter evaluation dataset
- * Datasets with fewer langs
- * Extract from Wikipedia
- + smoke test

Challenges

- * Linguistic challenges

- know next to nothing about 90% of the languages
- languages and scripts

<http://www.omniglot.com/writing/langalph.htm>

- language distributions

https://en.wikipedia.org/wiki/Languages_used_on_the_Internet#Content_languages_for_websites

- word segmentation?

- * Engineering challenges

NLP Evaluation


Intrinsic: use a gold-standard dataset and some metric to measure the system's performance directly.
(in-domain & out-of-domain)

Extrinsic: use a system as part of an upstream task(s) and measure performance change there.


Metrics

**“IF YOU CAN’T
MEASURE IT, YOU
CAN’T MANAGE IT”**

PETER DRUCKER



Dev-Test Split



f1 et al.

		Condition (as determined by "Gold standard")			
Total population		Condition positive	Condition negative	Prevalence = $\frac{\Sigma \text{ Condition positive}}{\Sigma \text{ Total population}}$	
Test outcome	Test outcome positive	True positive	False positive (Type I error)	Positive predictive value (PPV, Precision) = $\frac{\Sigma \text{ True positive}}{\Sigma \text{ Test outcome positive}}$	False discovery rate (FDR) = $\frac{\Sigma \text{ False positive}}{\Sigma \text{ Test outcome positive}}$
	Test outcome negative	False negative (Type II error)	True negative	False omission rate (FOR) = $\frac{\Sigma \text{ False negative}}{\Sigma \text{ Test outcome negative}}$	Negative predictive value (NPV) = $\frac{\Sigma \text{ True negative}}{\Sigma \text{ Test outcome negative}}$
Positive likelihood ratio (LR+) = TPR/FPR		True positive rate (TPR, Sensitivity, Recall) = $\frac{\Sigma \text{ True positive}}{\Sigma \text{ Condition positive}}$	False positive rate (FPR, Fall-out) = $\frac{\Sigma \text{ False positive}}{\Sigma \text{ Condition negative}}$	Accuracy (ACC) = $\frac{\Sigma \text{ True positive} + \Sigma \text{ True negative}}{\Sigma \text{ Total population}}$	
Negative likelihood ratio (LR-) = FNR/TNR		False negative rate (FNR) = $\frac{\Sigma \text{ False negative}}{\Sigma \text{ Condition positive}}$	True negative rate (TNR, Specificity, SPC) = $\frac{\Sigma \text{ True negative}}{\Sigma \text{ Condition negative}}$		
Diagnostic odds ratio (DOR) = LR+/LR-					


f1 et al.

$$F_1 = \frac{2}{\frac{1}{\text{recall}} + \frac{1}{\text{precision}}} = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

$$F_\beta = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}$$

$$G = \sqrt{\text{precision} \cdot \text{recall}}$$

ROC AUC



Confusion Matrix

Languages	English	German	French	Italian	Dutch	Spanish
English	9244	38	199	145	222	139
German	28	9514	67	29	325	27
French	20	52	9525	165	83	160
Italian	6	7	18	9822	16	134
Dutch	60	66	35	20	9800	19
Spanish	6	8	41	242	24	9679

AF: 1.00 |

DE: 1.00 |

EN: 1.00 |

ES: 0.94 | IT:0.06

NL: 0.85 | IT:0.03 CA:0.03 AF:0.03 DE:0.03 FR:0.03

RU: 0.82 | UK:0.12 EN:0.03 DA:0.02 FR:0.02

UK: 0.93 | TG:0.07

Total quality: 0.91

Evaluating “NLG”

- 1) Word Error Rate (WER)
- 2) BLEU, ROUGE, METEOR

BLEU Evaluation Metric

(Papineni et al, ACL-2002)

Reference (human) translation:

The U.S. island of Guam is maintaining a high state of alert after the Guam airport and its offices both received an e-mail from someone calling himself the Saudi Arabian Osama bin Laden and threatening a biological/chemical attack against public places such as the airport.

Machine translation:

The American [?] international airport and its the office all receives one calls self the sand Arab rich business [?] and so on electronic mail , which sends out ; The threat will be able after public place and so on the airport to start the biochemistry attack , [?] highly alerts after the maintenance.

- BLEU is a weighted geometric mean, with a brevity penalty factor added.
 - Note that it's precision-oriented
- BLEU4 formula
(counts n-grams up to length 4)

$$\exp (1.0 * \log p1 + \\ 0.5 * \log p2 + \\ 0.25 * \log p3 + \\ 0.125 * \log p4 - \\ \max(\text{words-in-reference} / \text{words-in-machine} - 1, 0))$$

p1 = 1-gram precision

P2 = 2-gram precision

P3 = 3-gram precision

P4 = 4-gram precision

Note: only works at corpus level (zeroes kill it);
there's a smoothed variant for sentence-level

Other Metrics

- 1) Cross-entropy
- 2) Perplexity
- 3) Custom: MaxMatch, Tribunal,...
- 4) Your own?

Baseline

Quality that can be achieved using some reasonable primitive approach (on the same data).

2 ways to measure improvement:

- * quality improvement
- * error reduction

Recent SNLI Example

<https://arxiv.org/pdf/1803.02324.pdf>

- majority class baseline: 0.34
- SOTA models: 0.87

Recent SNLI Example

<https://arxiv.org/pdf/1803.02324.pdf>

- majority class baseline: 0.34
- SOTA models: 0.87
- basic fasttext classifier: 0.67

	Entailment		Neutral		Contradiction	
SNLI	outdoors	2.8%	tall	0.7%	nobody	0.1%
	least	0.2%	first	0.6%	sleeping	3.2%
	instrument	0.5%	competition	0.7%	no	1.2%
	outside	8.0%	sad	0.5%	tv	0.4%
	animal	0.7%	favorite	0.4%	cat	1.3%
MNLI	some	1.6%	also	1.4%	never	5.0%
	yes	0.1%	because	4.1%	no	7.6%
	something	0.9%	popular	0.7%	nothing	1.4%
	sometimes	0.2%	many	2.2%	any	4.1%
	various	0.1%	most	1.8%	none	0.1%

Table 4: Top 5 words by $\text{PMI}(\text{word}, \text{class})$, along with the proportion of *class* training samples containing *word*. MultiNLI is abbreviated to MNLI.




Figure 1: The probability mass function of the hypothesis length in SNLI, by class.

State-of-the-Art (SOTA)

The highest publicly known
result on a well-established
dataset.

https://aclweb.org/aclwiki/State_of_the_art

GLUE Benchmark

The General Language Understanding Evaluation (GLUE) benchmark is a collection of resources for training, evaluating, and analyzing natural language understanding systems. GLUE consists of:


- * A benchmark of nine sentence- or sentence-pair language understanding tasks built on established existing datasets and selected to cover a diverse range of dataset sizes, text genres, and degrees of difficulty
- * A diagnostic dataset designed to evaluate and analyze model performance with respect to a wide range of linguistic phenomena found in natural language, and
- * A public leaderboard for tracking performance on the benchmark and a dashboard for visualizing the performance of models on the diagnostic set.

Evolution of NLProc Paradigms

1. Symbolical/rule-based
“The (Weak) Empire of Reason”
2. Counting-based
“The Empiricist Invasion”
3. Deep Learning-based
“The Revenge of the Spherical Cows”

<http://www.earningmyturns.org/2017/06/a-computational-linguistic-farce-in.html>

Rule-based Approach



English Stemming

(rule-based definition)

Porter stemmer:

```
(defun stem (word)
  (if (<= (length word) 2)
      word
      (-> word step1ab step1c step2 step3 step4 step5)))
```

```
step1ab: (-> word stem-s stem-ed/ing)
stem-s:  (when (and (ends-with "s" word)
                    (not (ends-with "ss" word))))
          (substr word (if (or (ends-with "sses" word)
                                (ends-with "ies" word))
                            -2 -1)))
step1c:  (when (and (ends-with "y" word)
                    (vowel-in-stem? word 1))
          (:= (last-char word) #\i))
```

...

<https://github.com/vseloved/cl-nlp/blob/master/src/lexics/porter.lisp>

English Tokenization (small problem)

* Simplest regex:

```
[^\s]+
```

* More advanced regex:

```
\w+|["'#$%&'*,.\/:;<=>?@^`~...\\(\)\{\}\[\|\|]———«»“”‘’-]
```

* Even more advanced regex:

```
[+-]?[0-9](?:[0-9,\.]*[0-9])?  
|[\w@](?:[\w'`@-][\w']|[\w'][\w@'`-])*[\w']?  
|["#$%&*,./:;<=>@^`~...\\(\)\{\}\[\|\|]———«»“”‘’']  
|[\.!?]+  
|-+
```

Add post-processing:

* concatenate abbreviations and decimals

* split contractions with regexes

- 2-character abbreviations regex:

```
i['```]m|(?:s?he|it)['```]s|(?:i|you|s?he|we|they)['```]d$
```

- 3-character abbreviations regex:

```
(?:i|you|s?he|we|they)['``'](?:ll|[vr]e)|n['```]t$
```

https://github.com/lang-uk/tokenize-uk/blob/master/tokenize_uk/tokenize_uk.py

Error Correction (inherently rule-based)

LanguageTool:

```
<category name="Стиль" id="STYLE" type="style">
  <rulegroup id="BILSH_WITH_ADJ" name="Більш з прикметниками">
    <rule>
      <pattern>
        <token>більш</token>
        <token postag_regexp="yes" postag="ad[jv]:.*compc.*">
          <exception>менш</exception>
        </token>
      </pattern>
      <message>Після «більш» не може стояти вища форма прикметника</message>
      <!-- TODO: when we can bind comparative forms togher again
        <suggestion><match no="2"/></suggestion>
        <suggestion><match no="1"/> <match no="2" postag_regexp="yes"
postag="(.*):compc(.*)" postag_replace="$1:compb$2"/></suggestion>
        <example correction="Світліший|Більш світлий"><marker>Більш
світліший</marker>.</example>
        -->
        <example correction=""><marker>Більш світліший</marker>.</example>
        <example>все закінчилось більш менш</example>
      </rule>
```

...

<https://github.com/languagetool-org/languagetool/blob/master/languagetool-language-modules/uk/src/main/resources/org/languagetool/rules/uk/>

Fighting Spam

(just a bad idea :)

SpamAssassin:

```
# bodyn
# example: postacie greet! Chcesz porozmawiac  Co prawda tu rzadko bywam,
zwykle pisze tu - http://hanna.3xa.info
uri      __LOCAL_LINK_INFO    /http:\\\\\\w{3,8}\\w\\w\\w\\.info/
header   __LOCAL_FROM_02     From =~ /\@o2\.pl/
meta     LOCAL_LINK_INFO     __LOCAL_LINK_INFO && __LOCAL_FROM_02
describe LOCAL_LINK_INFO     Link postaci http://cos.cos.info i z o2.pl
score    LOCAL_LINK_INFO     5
```

<https://wiki.apache.org/spamassassin/CustomRulesets>

(...which continues being
reimplemented)

Facebook antispam (using HAXL):

```
fpSpammer :: Haxl Bool
fpSpammer =
  talkingAboutFP .&&
  numFriends .> 100 .&&
  friendsLikeCPlusPlus
where
  talkingAboutFP =
    strContains "Functional Programming" <$> postContent

  friendsLikeCPlusPlus = do
    friends <- getFriends
    cppFriends <- filterM likesCPlusPlus friends
    return (length cppFriends >= length friends `div` 2)
```

<http://multicore.doc.ic.ac.uk/iPr0gram/slides/2015-2016/Marlow-fighting-spam.pdf>

https://petrimazepa.com/m_li_ili_kak_algoritm_facebook_blokiruet_polzovat_elei_na_osnovanii_stop_slov

Dialog Systems (inherently scripted)

ELIZA:

```
(defparameter *eliza-rules*  
'((((?* ?x) hello (* ?y))  
  (How do you do. Please state your problem.))  
  (((?* ?x) I want (* ?y))  
    (What would it mean if you got ?y)  
    (Why do you want ?y) (Suppose you got ?y soon))  
  (((?* ?x) if (* ?y))  
    (Do you really think its likely that ?y) (Do you wish that ?y)  
    (What do you think about ?y) (Really-- if ?y))  
  (((?* ?x) no (* ?y))  
    (Why not?) (You are being a bit negative)  
    (Are you saying "NO" just to be negative?))  
  (((?* ?x) I was (* ?y))  
    (Were you really?) (Perhaps I already knew you were ?y)  
    (Why do you tell me you were ?y now?))  
  (((?* ?x) I feel (* ?y))  
    (Do you often feel ?y ?))  
  (((?* ?x) I felt (* ?y))  
    (What other feelings do you have?))))
```

<https://norvig.com/paip/eliza1.lisp>

...and if you thought rules are dead

Wit.ai:

To make a bot, there are two schools: rules or machine learning. (Everybody claims rules are bad and their bot is powered by AI, but when you really look under the hood, the core is often imperative.)



Machine learning is of course more desirable, but the problem is the training dataset. Training a Wit intent with a dozen examples works well, and it's easy to leverage the community to get more examples. But in order to entirely learn the business logic of a bot of medium complexity, we would need many, many thousands of example conversations.

Rules (or any kind of imperative approach, including plain script/program) are kind of the opposite. The good thing with rules is, you can have a demo working after you write two rules. As long as you follow the script carefully, your bot will work and your audience will be impressed. But as you discover new “paths” in the dialog, you'll add more and more rules, until one day everything collapses. You're doomed by the curse of combinatorics. Any new rule conflicts with old rules that you totally forgot the reason for. Your bot cannot improve anymore.

When you create your bot, you just start with a few stories that describe the most probable conversations paths. At this stage, Bot Engine will build a machine learning model that deliberately overfits the stories dataset. Practically, it means that stories will behave almost like rules.

<https://medium.com/wit-ai/bot-engine-26af22d37fd6>


 Create story


I want to go see **Straight Outta Compton**


movie_name


Straight Outta Compton

 Add a new entity

say (Where would you like to go watch **[movie_name]**?)




 Add variable




ACME theatre in NY


wit/location


ACME theatre in NY

 Add a new entity

say (What time and how many tickets?)



 Add variable




4 tickets today at 7pm please

wit/datetime


4/8/2016, 7:00:00 PM

wit/number


4

 Add a new entity


book_tickets (context, entities)





Doesn't produce context

 Add context field


say (You're all set!)




 Add variable



User says...



Bot says...



Bot executes...

More Rule-based Examples

- * Rule-based MT (Systran)
- * Various NLG systems
- * Custom parsing (Sparsen)

Event Detection at Ravenpack

Augmented Match Example 1

www.ravenpack.com

"MEXICO: Consumer Confidence Index Rises 1% In September On Monthly Basis"

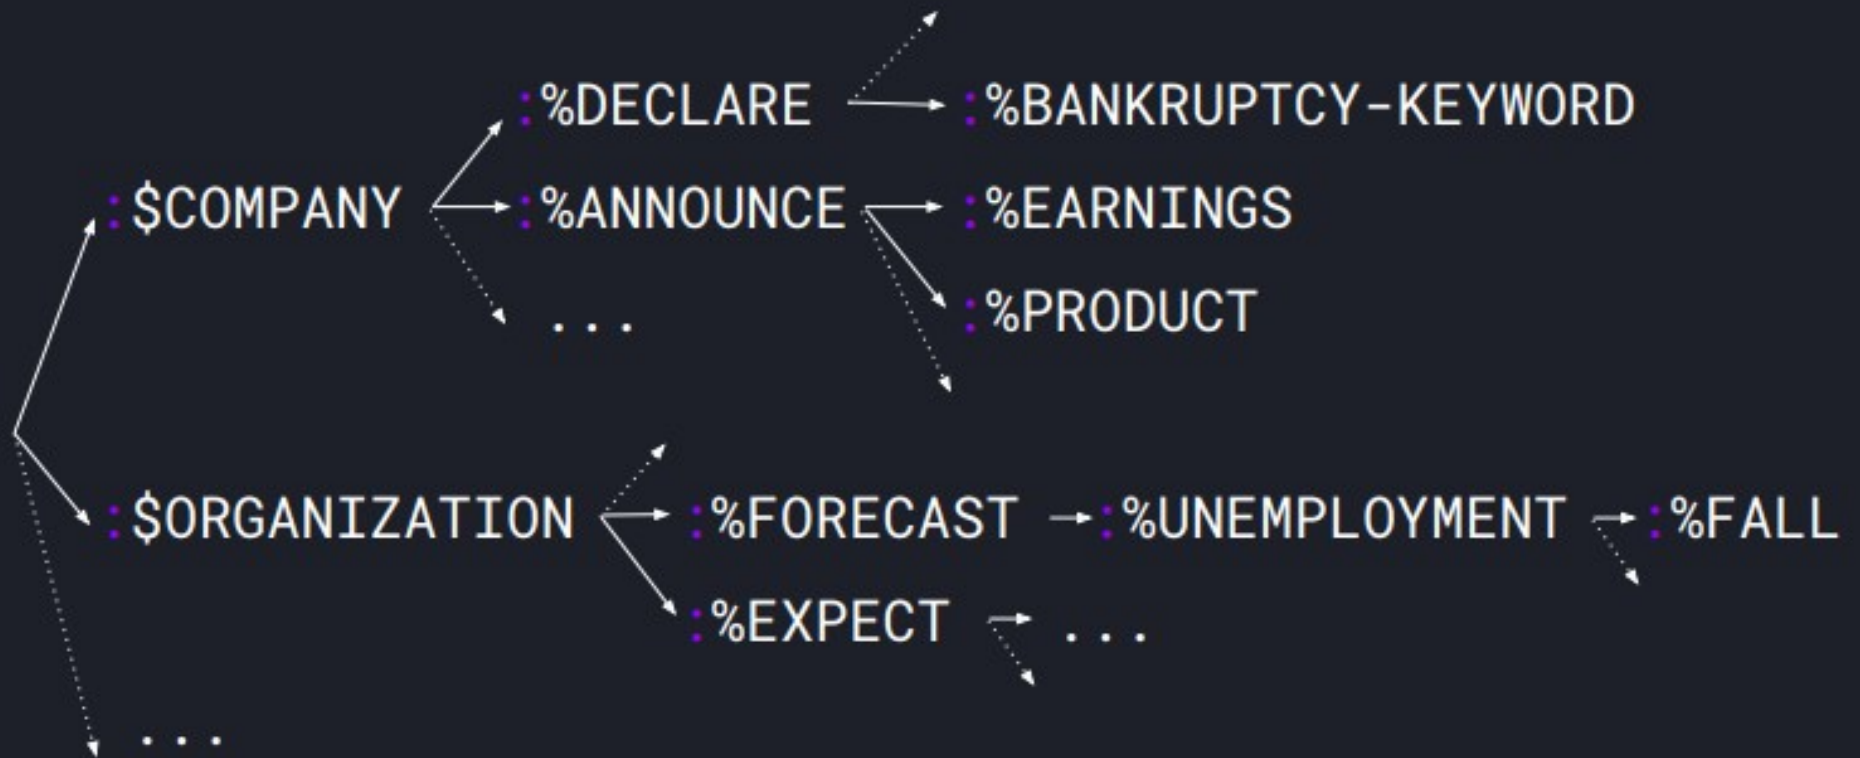
(:\$PLACE :%CONSUMER-CONFIDENCE :%RISE :\$PERCENTAGE :%PREPOSITION-TIME :\$PERIOD)

\$PLACE	%CONSUMER-CONFIDENCE	%RISE	\$PERCENTAGE	%PREPOSITION-TIME	\$PERIOD(*)	\$PERIOD
MEXICO	CONSUMER CONFIDENCE INDEX	RISES	1	IN	SEPTEMBER	ON MONTHLY BASIS
1.	consumer-confidence-up	42125631020B6113DF09246655383EE9 - Mexico (ESS: 0.53)				
2.	consumer-confidence-up-actual					
3.	consumer-confidence-up-previous					
4.	consumer-confidence-up-percentage	1	TIMESPAN: 2017-09-01 00:00:00 -> 2017-10-01 00:00:00 (Period: -----)			
5.	consumer-confidence-up-period					
6.	consumer-confidence-up-rater					
7.	consumer-confidence-up-date					
8.	consumer-confidence-up-method	METH: MOM				
9.	consumer-confidence-up-sentiment-modifier					
10.	consumer-confidence-up-sentiment					
11.	consumer-confidence-up-opinion					
12.	consumer-confidence-up-previous-period					

[https://european-lisp-symposium.org/static/2018/corn
ez.pdf](https://european-lisp-symposium.org/static/2018/corn ez.pdf)

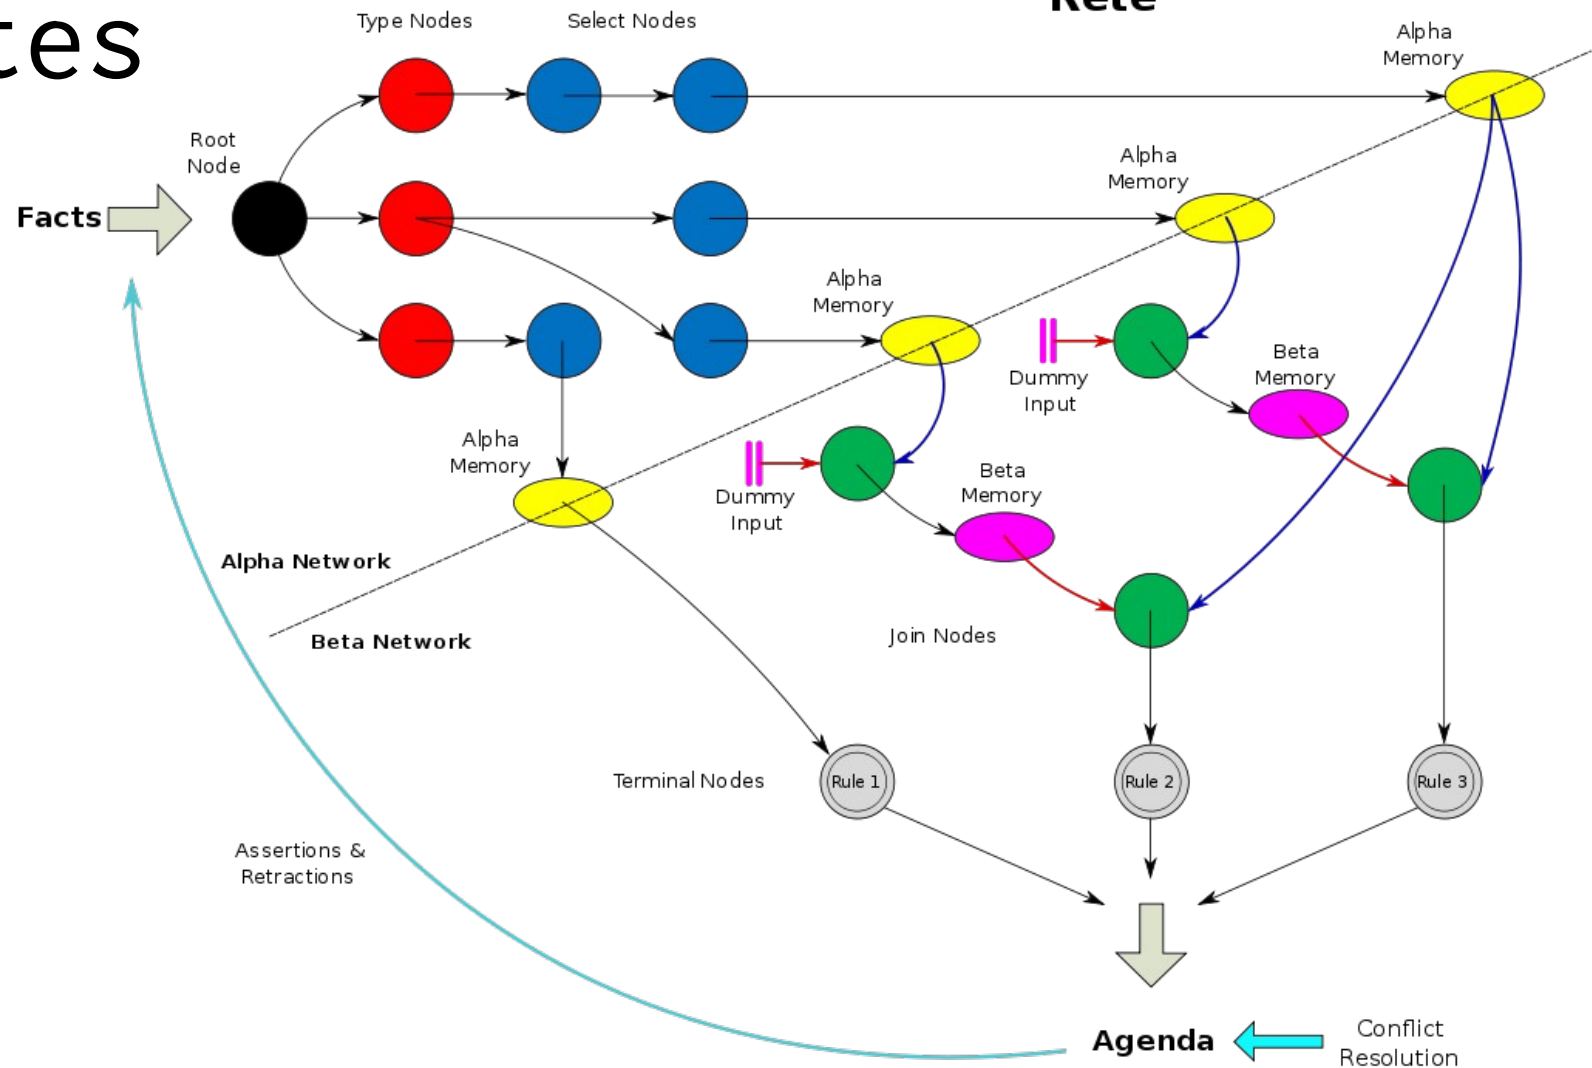
A Template Trie

WW



Rete Algorithm

Efficient matching of many
templates



Languages for Rules


- * regexes

Languages for Rules

- * regexes

Languages for Rules

- * regexes
- * XML, JSON



Languages for Rules

- * regexes
- * XML, JSON
- * external DSLs (example: JBOSS drules)

Listing 5.7 A simple set of rules for email spam filtering

```
package demo;
import iweb2.ch5.classification.data.Email;
import iweb2.ch5.classification.rules.ClassificationResult;

global ClassificationResult classificationResult;

rule "Tests for viagra in subject"
when
    Email( $s : subject )
    eval( classificationResult.isSimilar($s, "viagra" ) )
then
    classificationResult.setSpamEmail(true);
end

rule "Tests for 'drugs' in subject"
when
    Email( $s : subject )
    eval( classificationResult.isSimilar($s, "drugs" ) )
then
    classificationResult.setSpamEmail(true);
end
```

**Rule for identifying
"Viagra" in email
subject**

**Rule for identifying
"drugs" in email
subject**

Languages for Rules

- * regexes
- * XML, JSON
- * external DSLs
- * internal DSLs

```
(match-html source
  '(>> article
    (aside (>> a ($ user))
      (>> li (strong "Native Tongue:") ($ lang)))
    (div |...| (>> (div :data-role "commentContent"
      ($ text) (span) |...|)))
    !!!))
```

<https://github.com/vseloved/crawlik>

Rules Pros&Cons

- + compact & fast
- + full control
- + arbitrary recall
- + iterative
- + best interpretability
- + perfectly accommodate domain experts
- precision ceiling
- non-optimal weights
- require a lot of human labor
- hard to interpret/calibrate score

Hybrid Approach

“Rule-based” framework that incorporates Machine Learning models.


- + control
- + best of both worlds
- more complex
- not sexy

Counteracting Bias

- * Rule-based post-processing
- * Mix multiple datasets
+ special feature-selection
- * Train multiple models on
different datasets and
create an ensemble
- * Use domain adaptation
techniques

Productization

Delivering your NLP application to the users



Product variants

- * end-user product
(web, mobile, desktop)
- * internal feature
- * API
- * library
- * scientific paper

Requirements

- * speed (processing, startup)
- * memory usage
- * storage
- * interoperability
- * ease of use
- * ease of update

Storage Optimization

WILD Initial model size ~ 1G
(compared to megabytes for CLD,
target: 10-20 MB)


Ways to reduce:

- partly rule-based
- model pruning
- clever algorithms: perfect hash tables, Huffman coding
- model compression

<https://arxiv.org/abs/1612.03651>

Model Formats

* pickle & co.



Model Formats

- * pickle & co.
- * JSON, ProtoBuf, ...

Model Formats

- * pickle & co.
- * JSON, ProtoBuf, ...
- * HDF5

Model Formats

- * pickle & co.
- * JSON, ProtoBuf, ...
- * HDF5
- * Custom gzipped text-based

Adaptation

(manual vs automatic)

- * for RB systems: add more rules :D
- * for ML systems: online learning algorithms
- * for DL systems: transfer Learning
- * Lambda architecture

The Pipeline

“Getting a data pipeline into production for the first time entails dealing with many different moving parts – in these cases, spend more time thinking about the pipeline itself, and how it will enable experimentation, and less about its first algorithm.”

https://medium.com/@neal_lathia/five-lessons-from-building-machine-learning-systems-d703162846ad

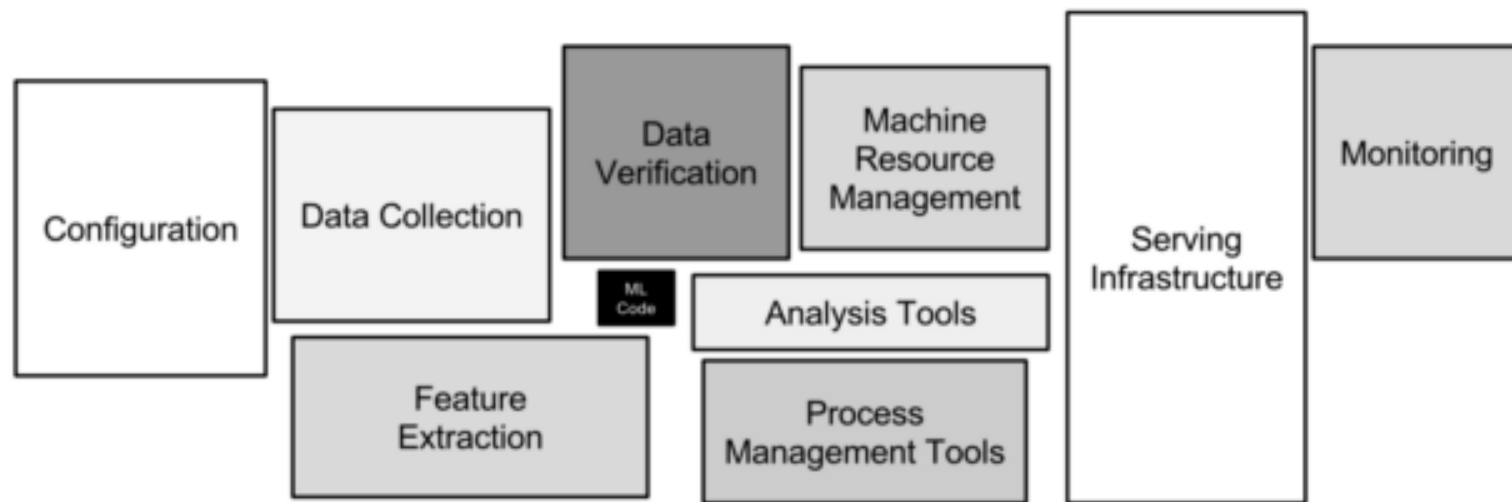


Figure 1: Only a small fraction of real-world ML systems is composed of the ML code, as shown by the small black box in the middle. The required surrounding infrastructure is vast and complex.

THE DATA SCIENCE **HIERARCHY OF NEEDS**

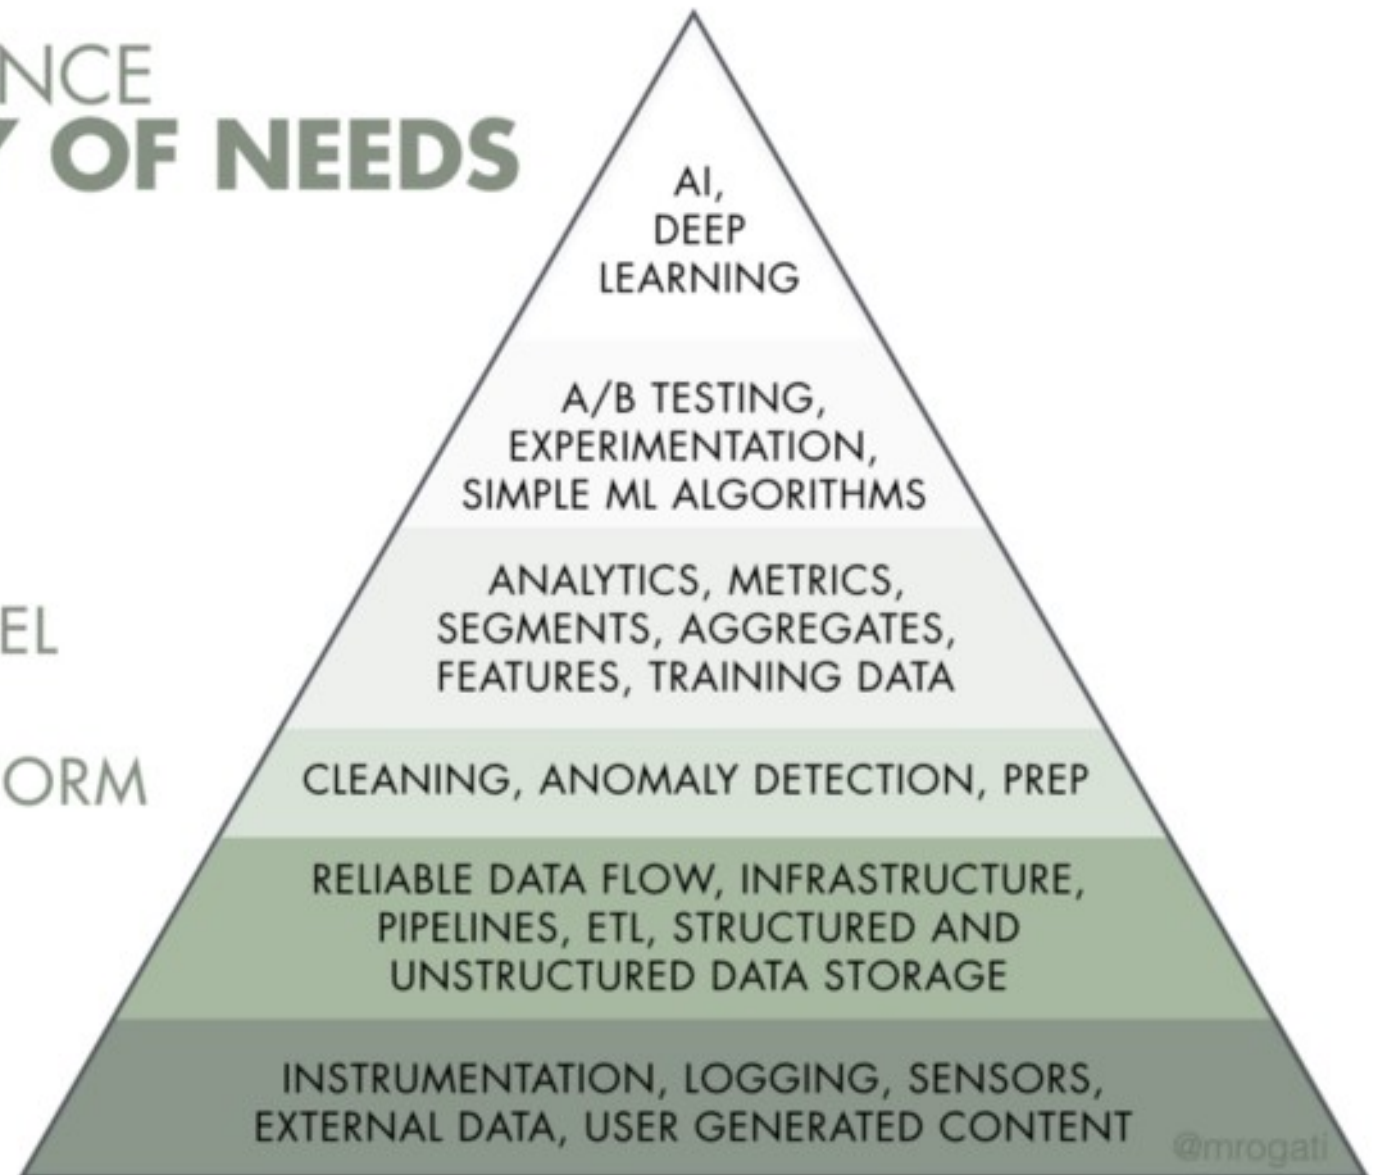
LEARN/OPTIMIZE

AGGREGATE/LABEL

EXPLORE/TRANSFORM

MOVE/STORE

COLLECT



<https://hackernoon.com/the-ai-hierarchy-of-needs-18f111fcc007>

Read More

- <http://people.cs.umass.edu/~brenocon/inlp2015/15-eval.pdf>
- <https://ehudreiter.com/2017/05/03/metrics-nlg-evaluation/>
- <https://nlpers.blogspot.com/2014/11/the-myth-of-strong-baseline.html>
- <http://kavita-ganesan.com/what-is-rouge-and-how-it-works-for-evaluation-of-summaries>
- <https://www.slideshare.net/frandzi/native-language-identification-brief-review-to-the-state-of-the-art>
- <https://www.slideshare.net/grammarly/grammarly-a-inlp-club-1-domain-and-social-bias-in-nlp-case-study-in-language-identification-tim-baldwin-80252288>